

COMPARISON OF N-STEP ALGORITHMS

Jonathan Campbell


COMP-767

February 17, 2017

OVERVIEW

- Comparison of off-policy n-step algorithms to estimate Q
 - Off-policy n-step Expected Sarsa
 - n-step Tree Backup
 - Off-policy n-step $Q(\sigma)$
- Values for α and σ will be compared.

N-STEP EXPECTED SARSA



Use n-step return in update rule:

$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha \left[G_t^{(n)} - Q_{t+n-1}(S_t, A_t) \right]$$

Return is sum of rewards plus expectation at last timestep:

$$G_t^{(n)} \doteq R_{t+1} + \dots + \gamma^{n-1} R_{t+n} + \gamma^n \sum_a \pi(a|S_{t+n}) Q_{t+n-1}(S_{t+n}, a)$$

OFF-POLICY N-STEP EXPECTED SARSA

Weight updates by IS ratio ρ :


$$Q_{t+n}(S_t, A_t) \doteq Q_{t+n-1}(S_t, A_t) + \alpha \rho_{t+1}^{t+n-1} \left[G_t^{(n)} - Q_{t+n-1}(S_t, A_t) \right]$$

ρ is relative prob. under target/behavior policy of taking the sequence of actions:


$$\rho_t^{t+n} \doteq \prod_{k=t}^{\min(t+n-1, T-1)} \frac{\pi(A_k | S_k)}{\mu(A_k | S_k)}$$

Trajectories could be discarded if one action has zero prob. under target policy.

Also: high IS ratios could cause high variance in Q-values, requiring lower learning rate.



N-STEP TREE BACKUP



Off-policy learning without importance sampling.

N-step return considers all possible actions from each state (action q-val multiplied by action prob. under π).

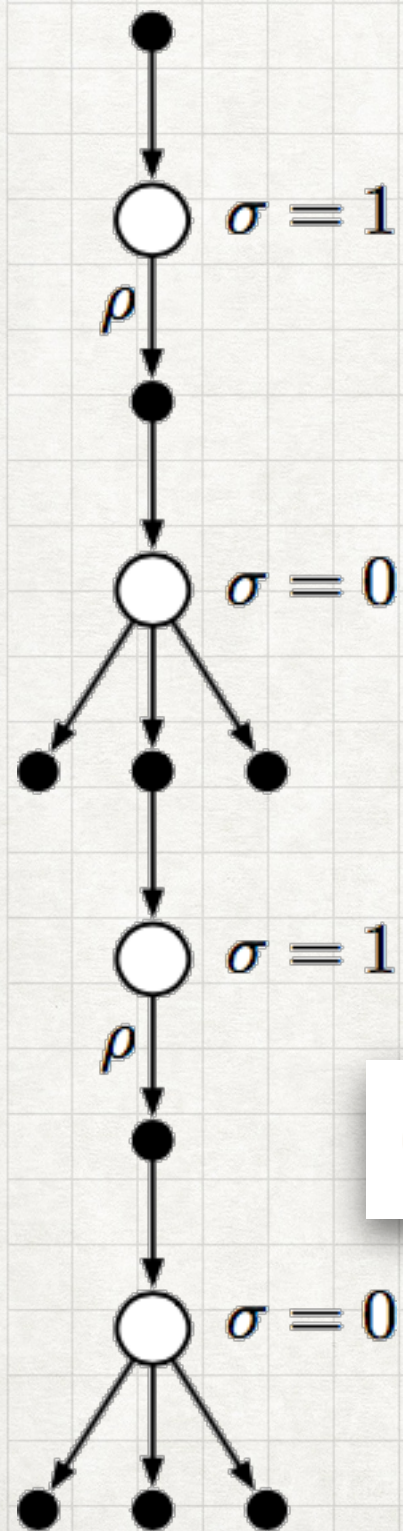
The action taken in a particular state takes into account the reward observed and actions from the next state.

2-step case:

$$G_t^{(2)} = R_{t+1} + \gamma V_{t+1} + \gamma \pi(A_{t+1}|S_{t+1}) [R_{t+2} + \gamma V_{t+2} - Q_t(S_{t+1}, A_{t+1})]$$

where $V_t = \sum_a \pi(a|S_t) Q_{t-1}(S_t, a)$

N-STEP Q(SIGMA)



Generalization of expected Sarsa and tree-backup algorithm.

Mix Sarsa update and tree backup update, depending on σ .

$$G_t^{(n)} \doteq Q_{t-1}(S_t, A_t) + \sum_{k=t}^{\min(t+n-1, T-1)} \delta_k \prod_{i=t+1}^k \gamma [(1 - \sigma_i) \pi(A_i | S_i) + \sigma_i]$$


$$\delta_t \doteq R_{t+1} + \gamma [\sigma_{t+1} Q_t(S_{t+1}, A_{t+1}) + (1 - \sigma_{t+1}) V_{t+1}] - Q_{t-1}(S_t, A_t)$$

(Also have to modify IS ratio for off-policy case.)

IMPLEMENTATION

- Uses Gridworld RL framework from UC Berkeley AI course
 - <http://ai.berkeley.edu/reinforcement.html>
- Extension to windy gridworld:

```
def getWindyGrid():  
    grid = [[' ', ' ', ' ', ' ', '^', '^', '^', '^', ' ', ' '],  
            [' ', ' ', ' ', '^', '^', '^', '^', ' ', ' '],  
            [' ', ' ', ' ', '^', '^', '^', '^', ' ', ' '],  
            [' ', ' ', ' ', '^', '^', '^', 1, ' ', ' '],  
            [' ', ' ', ' ', '^', '^', '^', '^', ' ', ' '],  
            [' ', ' ', '^', '^', '^', '^', '^', '^', ' ', ' '],  
            ['S', ' ', ' ', '^', '^', '^', '^', ' ', ' ']]  
    return Gridworld(grid)
```

CURRENT Q-VALUES


SIGMA PARAMETER

- Sigma parameter controls degree of sampling.
 - $\sigma = 0$: Tree-backup algorithm
 - $\sigma = 1$: Sarsa
- Values
 - Fixed
 - Linear/exponential decay
 - Function of current state and/or action


METHOD

- For each alpha/sigma value:
 - For 100 iterations:
 - Run n-step agent ($n=10$) on windy gridworld for 100 episodes and record total return of last episode.
 - Average the 100 total returns.
- For all algs., behaviour policy is ϵ -greedy.

SIGMA COMPARISON



N-STEP ALGORITHM COMPARISON



CODE

- Available at course GitHub repo:

<https://github.com/rllabmcgill/rlcourse-february-17-campbelljc>